

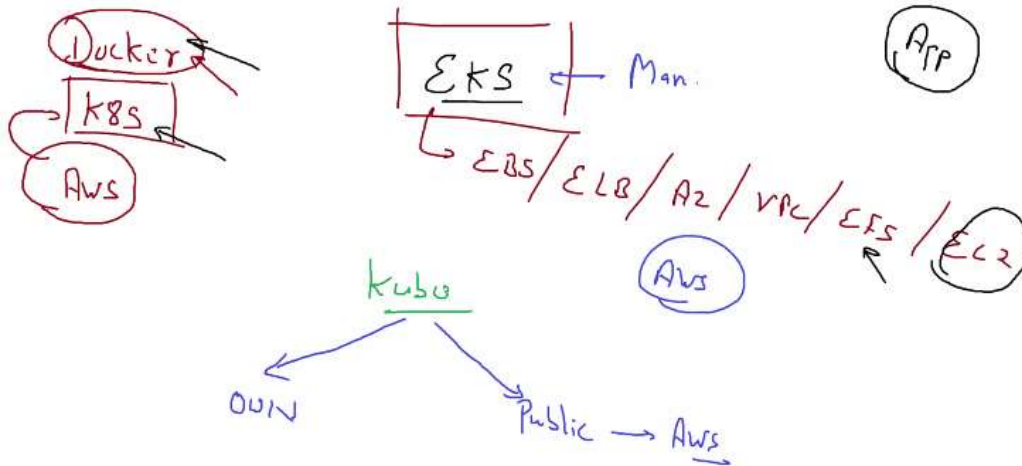
Day 1

02 July 2020 15:08

Actually we have the Docker, K8s, and AWS which are great services we have And might have a lot of uses and use case where they are used

But we have EKS the Elastic Kubernetes as Service which is the K8s managed on the AWS Cloud

- Which gives us the all the power and features of the (Public) Cloud
- And will be managed by the AWS people (the developers)



- And it is the form of Managed K8s on the Cloud (as the managing of anything is very important)
 - o Like for managing the Docker we use the K8s
 - o We are using the EKS which is the Managed K8s

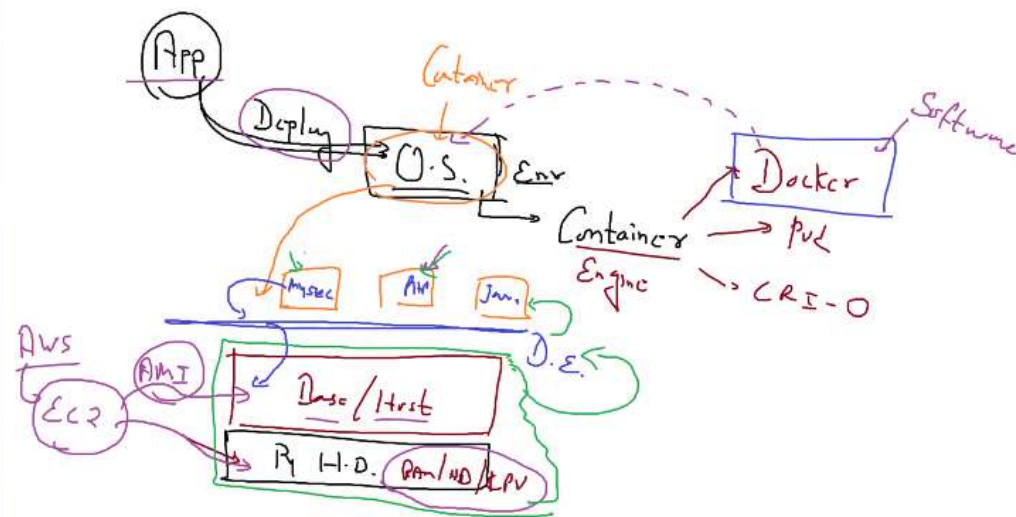
And also in this training we will be using and integrating the EKS with the multiple services

And this is a mature one better than what we can do in our own setup of K8s

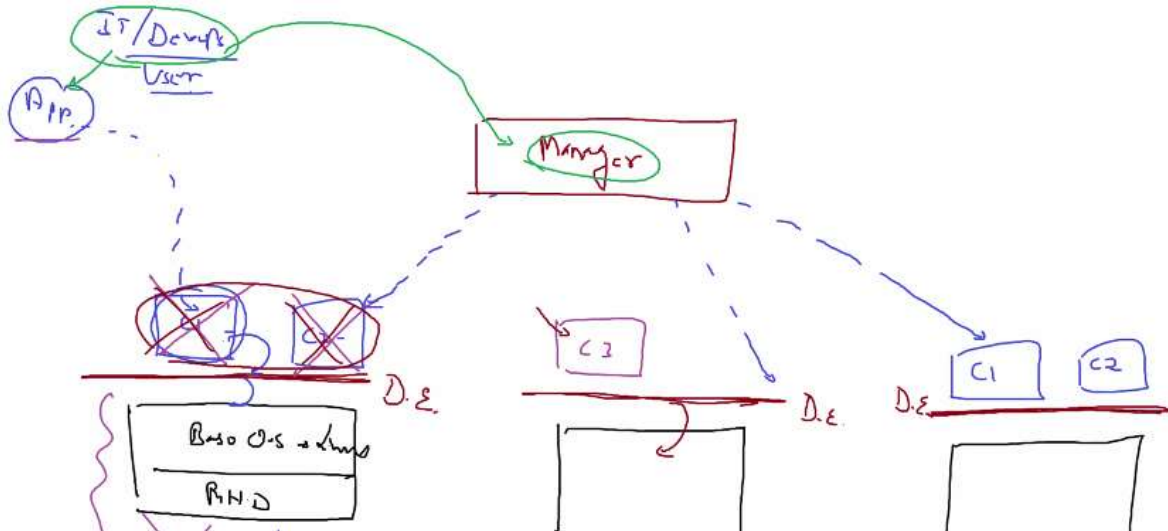
- even if we are thops in K8s,
- we can't rely on the Base system,
- and if we want our base resources or Base System to be mature
- (each and every part of it) then we have the service EKS

Some Basic:

- And for anything we need to have an OS
 - o As we might need to run an App or launch or run a program
- And for this we might use the Containers (which is most used now-a-days as within 1 sec we get the OS and also the isolation)



- And one of the Container technology is the Docker
- And to launch a container we need the Base Hardware on top of these we can launch it
- But for that we need to have those
- And these can be bought or we can use the CC (like AWS) for the resources
- And for the resources in AWS we have the service called as EC2 using which we can launch an OS or instance
- And then in the Docker there might be issues where the Container might deleted or closed and for this we run the same container in replicas called as Clustering
 - o Where we have the multiple Nodes and in those we have multiple containers to be launched
- But we might have the problem of managing cluster and this is done using the K8s
 - o Means like to relaunch it if it has deleted
 - o This is called as Fault Tolerance
 - o And many others



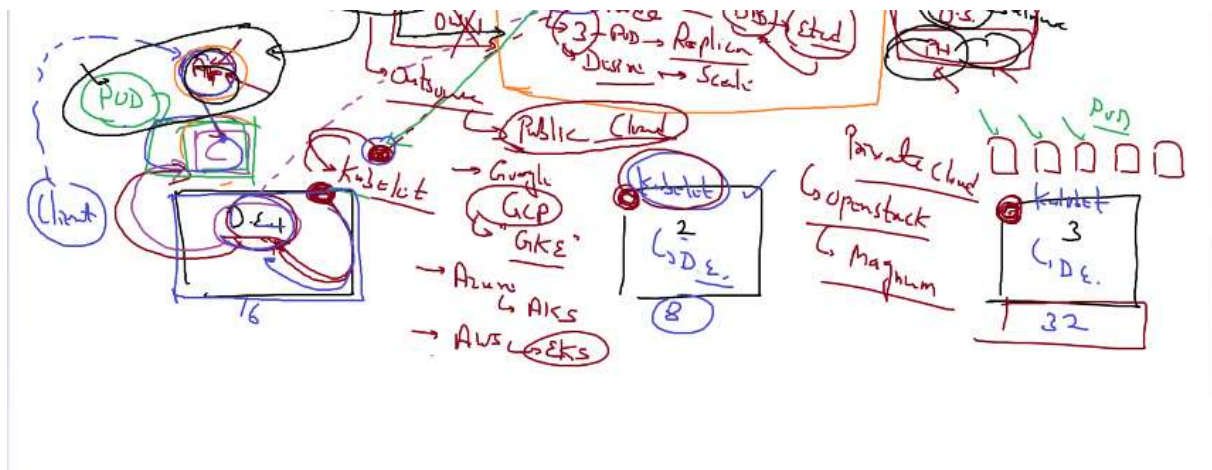
- So for all these we need a program (to make this automatic) which will manage the nodes, containers and others etc.,
- And this program is called as Container Orchestration Engine
- And one of the many which is famous is the K8s
- And the Managing node is called as master and the others are called as Slaves and these are called Master Slave Architecture
- And for connecting to the master and also to use it we need a program to which we can connect and this program here in K8s is called as API Server
- And the API server will be contacting to the Scheduler
- And the Scheduler is the one that will decide on which Node to use or launch the containers
 - o And the one who will launch the containers is the Docker but the Scheduler is the one to decide which node to be used
- And also the K8s or Scheduler don't know how to contact to the Docker so we need a program that will contact to the Docker and this is called as Kubelet (here)

And here the Company is the CC companies like

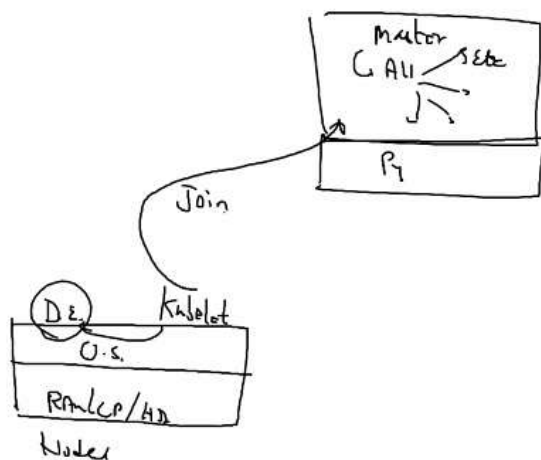
- AWS
- Google
- Azure (Microsoft)

And these have this service provided with their own name

In which they will also create the master Node and also



Also apart from the above we also need to even create and manage the entire cluster



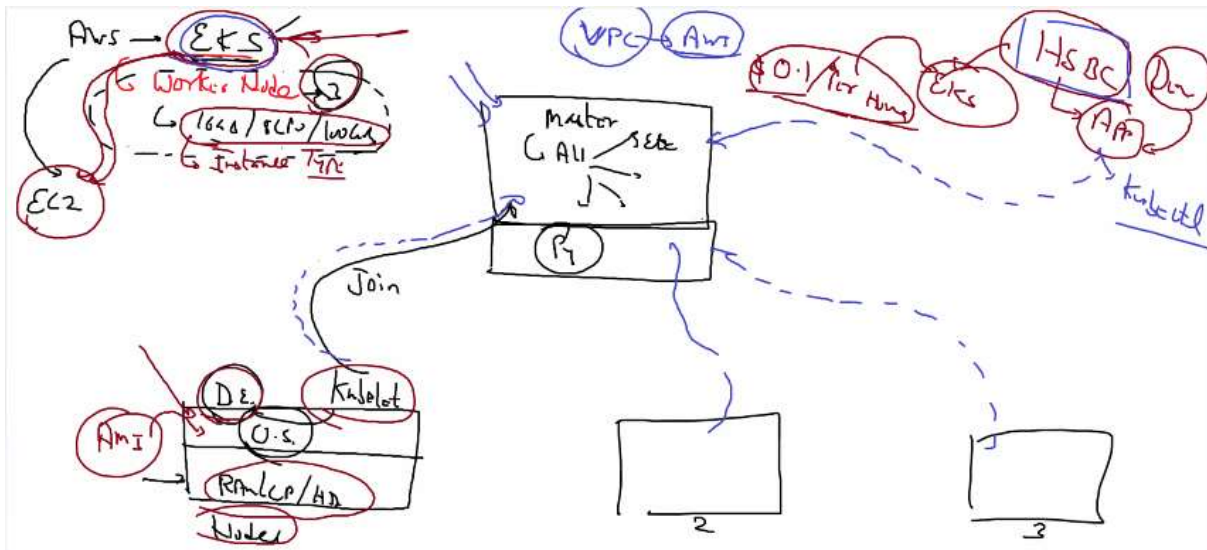
And all this we have the service in AWS cloud as the EKS service

- In which we can tell the no. of worker nodes we need
- The configuration we can tell what is the RAM and other resources we want for the slaves

And the charge for this cluster creation is the 0.1\$/ hour

And this is only for the Cluster creation and management

And for the other resources we are using in AWS they will be charged differently



And the beauty is that in one click the entire this is launched

And also this is cheap and also we don't need to manage the resources we are using and also it is launched using the best practices

And for the company this is the very big relief or a benefit to use this kind of service

And in the Industry the best and most important thing is the APP

And for them all this managing will be done with ease and in one click by the EKS

And for this we need to have the AWS account

aws.amazon.com/eks/pricing/

Amazon EKS Pricing - Managed Kubernetes Service

You pay \$0.10 per hour for each Amazon EKS cluster that you create. You can use a single Amazon EKS cluster to run multiple applications by taking advantage of Kubernetes namespaces and IAM security policies. ... See detailed pricing information on the Amazon EKS pricing page.

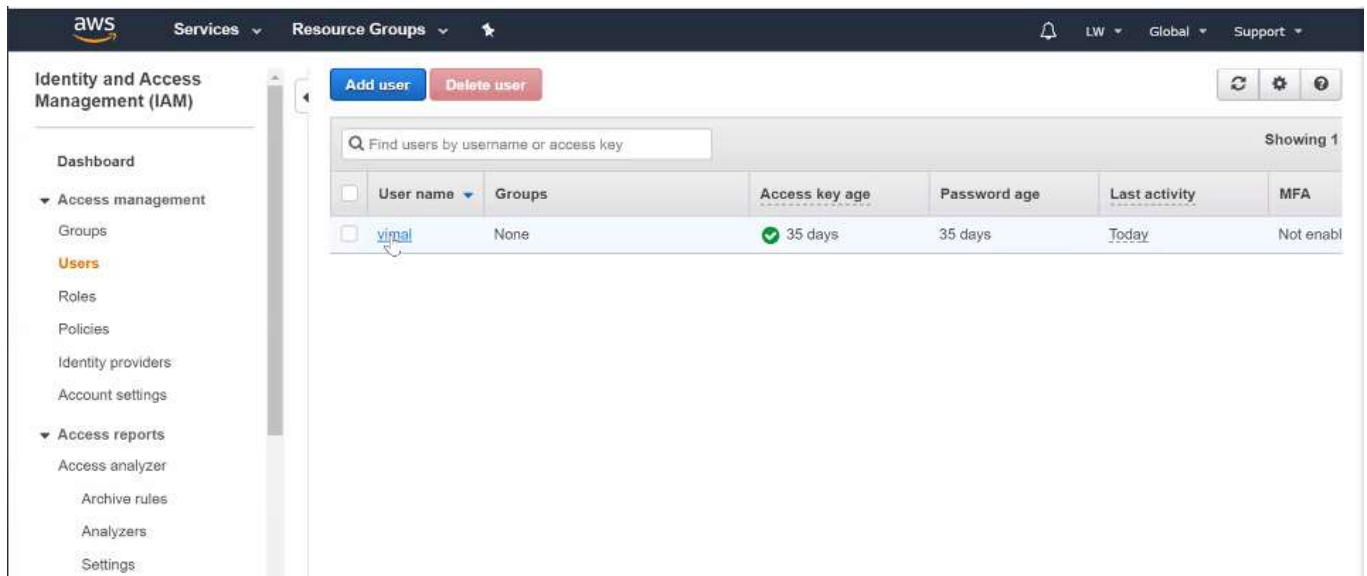
You visited this page on 3/7/20.

And they will charge for the per hour and also for per second which they have changed and this is a benefit

And to use the AWS we have many ways to use it

- Web UI
- CLI

And for the CLI we need to install the AWS CLI and check for the command if it does not work then we have to set the path then it works



And for using we need to login and for this we need the

- Username - Access Key
- Password - Secret Key

And for this we have or can create an account

And also while creating we need to specify the Role or Power we want to give it and for this we have to set the Admin for all the power if we want to give

```
C:\Users\Vimal Daga>aws configure
AWS Access Key ID [*****EIGQ]:
AWS Secret Access Key [*****92ua]:
Default region name [ap-south-1]:
Default output format [None]:
```

And to login we use the aws configure and then it will ask the username password default region and format

And then we have logged in perfectly

```
al Daga>aws eks list-clusters
```

And this is used to check the cluster we have (EKS)

```
C:\Users\chand>aws eks list-clusters
{
  "clusters": []
}
```

And to create the cluster we can use create-cluster

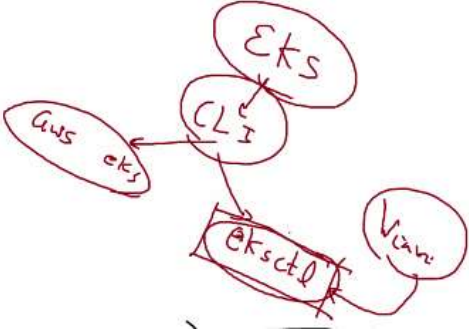
But the problem is that we can't do much customization with this command
As we don't have some much options for that

And the EKS is such a great service many companies might be hard to generate such customization

And from the CLI mode we have 2 ways now

- One is AWS typical aws eks command

- Or we have some other commands from some other companies (Veave here) is called as EKSCTL (EKS control)
 - o And this is a very powerful command
 - o That using this we can launch the cluster in one click with all the customizations we want



eksctl download

About 9,630 results (0.32 seconds)

docs.aws.amazon.com > eks > latest > userguide > gettin... ▾

Getting started with eksctl - Amazon EKS - AWS Documentation

(Optional) Verify the **downloaded** binary with the SHA-256 sum. **Download** the SHA-256 sum.
 curl -o ...

[Prerequisites](#) · [Install and configure kubectl](#) · [Create your Amazon EKS ...](#)

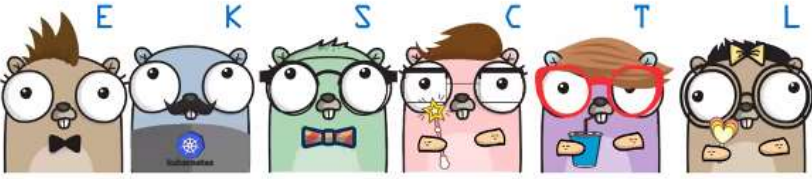
You visited this page on 3/7/20

eksctl - The official CLI for Amazon EKS

circled passing coverage: 38% go report A+

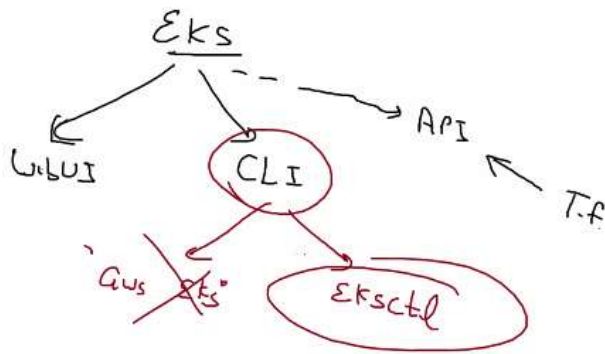
eksctl is a simple CLI tool for creating clusters on EKS - Amazon's new managed Kubernetes service for EC2. It is written in Go, and uses CloudFormation.

You can create a cluster in minutes with just one command - `eksctl create cluster` !



Need help? Join [Weave Community Slack](#).

And this software we can download it and then extract it and then set the path of the location where we have copied it



And this command is specialised for the EKS

```
C:\Users\Vimal Daga>eksctl version
0.21.0

C:\Users\Vimal Daga>eksctl get cluster
No clusters found
```

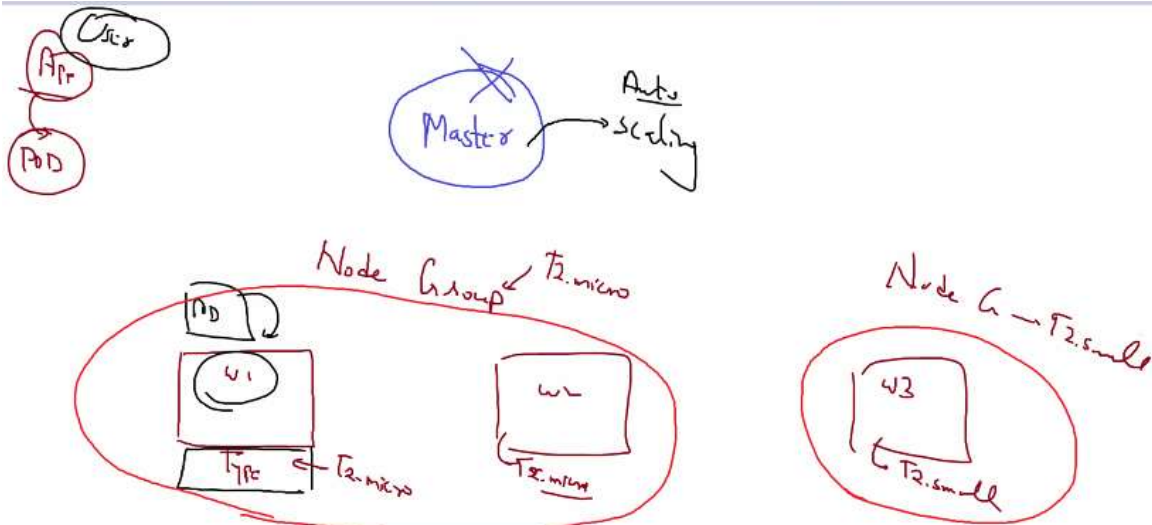
- And here to check the no. of cluster we have we use the `eks get cluster`

```
C:\Users\Vimal Daga>aws configure
AWS Access Key ID [*****EIGQ]:
AWS Secret Access Key [*****92ua]:
Default region name [ap-south-1]:
Default output format [None]:
```

And for this they use the login and password and they are taking it from the aws configure command only
And that's why we have to configure it before the install or using this

And this will also perform the auto scaling (the EKS service) and that is that if the cluster is not able to handle the load then they will automatically scale it

And also we might need the customization like which type of instance type to use



And this is called as Node Group

- And here if we are using two types of instance types then we are having 2 Node groups

And for all this we have to use some script and has to be created using the YAML language (it is a declarative language)


```
nodegroups:
ng1:
  2
  t2.micro
ng2:
  1
  t2.small
```

As per the above we need the 2 node group and also we have to specify which service to use
And also the no. of Nodes of such we need and this is called as capacity

And also has to specify the region where we want to launch

```
kind: ClusterConfig
region: ap-south=1
nodegroups:
ng1:
  capacity: 2
  type: t2.micro
ng2:
  cap: 1
  type: t2.small
```

And also about what are we doing and this is called as the Kind

And this is not the right syntax but this is the information which we have to provide

And also we have to give name to the cluster we are launching and this is called as metadata
And the final syntax looks like the

```
kind: ClusterConfig
metadata:
  name: lwcluster
  region: ap-south=1
nodeGroups:
  - name: ng1
    desiredCapacity: 2
    instanceType: t2.micro
  - name: ng2
    desiredCapacity: 1
    instanceType: t2.small
```

And here it has to be ap-south-1

And here it will launch the cluster in one click as we have said it that the kind is ClusteConfig

And also we have to tell which version to use as this kind will have been come from this some type of program file
And this program file here is called as apiVersion

```
apiVersion: eksctl.io/v1alpha5
```

```

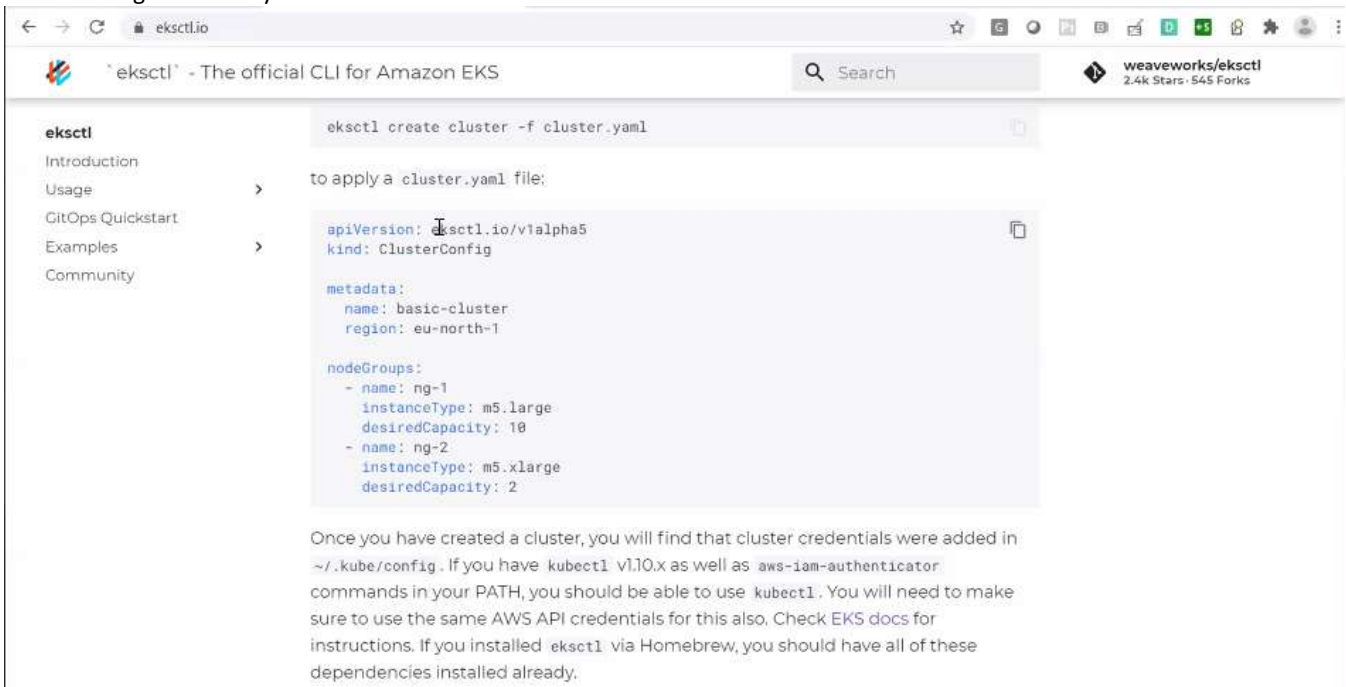
apiVersion: eksctl.io/v1alpha5
kind: ClusterConfig

metadata:
  name: lwcluster
  region: ap-south-1

nodeGroups:
  - name: ng1
    desiredCapacity: 2
    instanceType: t2.micro
  - name: ng2
    desiredCapacity: 1
    instanceType: t2.small

```

And we can get these keyword form the Official Documentation of the EKS



The screenshot shows the official website for eksctl, titled "eksctl - The official CLI for Amazon EKS". The page includes a search bar and a sidebar with links to Introduction, Usage, GitOps Quickstart, Examples, and Community. The main content area displays the command `eksctl create cluster -f cluster.yaml` and explains that it applies a `cluster.yaml` file. A sample configuration is provided:

```

apiVersion: eksctl.io/v1alpha5
kind: ClusterConfig

metadata:
  name: basic-cluster
  region: eu-north-1

nodeGroups:
  - name: ng-1
    instanceType: m5.large
    desiredCapacity: 10
  - name: ng-2
    instanceType: m5.xlarge
    desiredCapacity: 2

```

Below the configuration, text explains that after creating a cluster, credentials are added to `~/.kube/config`, and users can use `kubectl` if they have `kubectl v1.10.x` and `aws-iam-authenticator` installed. It also mentions that `eksctl` installed via Homebrew should have all dependencies already.

```

Directory of C:\Users\Vimal Daga\Desktop\eks_class_code
04-07-2020 19:54 <DIR> .
04-07-2020 19:54 <DIR> ..
04-07-2020 20:02      262 cluster.yaml
               1 File(s)      262 bytes
               2 Dir(s)  2,130,477,056 bytes free

C:\Users\Vimal Daga\Desktop\eks_class_code>eksctl get cluster
No clusters found

```

And when we run the command then they will create the cluster for us

```

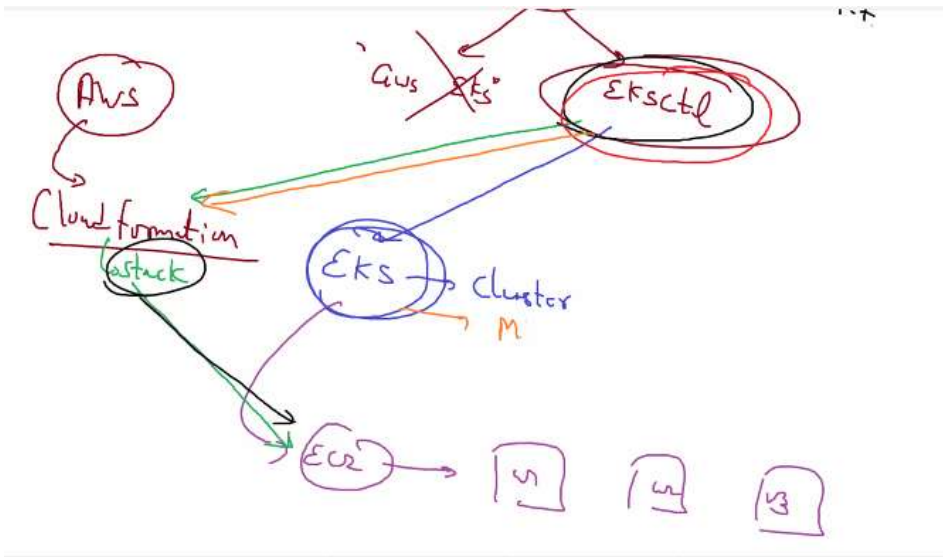
[0] nodegroup "ng2" will use "ami-073969767527f7306" [AmazonLinux2/1.16]
[0] using Kubernetes version 1.16
[0] creating EKS cluster "lwcluster" in "ap-south-1" region with un-managed nodes
[0] 2 nodegroups (ng1, ng2) were included (based on the include/exclude rules)
[0] will create a CloudFormation stack for cluster itself and 2 nodegroup stack(s)
[0] will create a CloudFormation stack for cluster itself and 0 managed nodegroup stack(s)
[0] if you encounter any issues, check CloudFormation console or try 'eksctl utils describe-
stacks --region=ap-south-1 --cluster=lwcluster'
[0] CloudWatch logging will not be enabled for cluster "lwcluster" in "ap-south-1"
[0] you can enable it with 'eksctl utils update-cluster-logging --region=ap-south-1 --cluste
r=lwcluster'
[0] Kubernetes API endpoint access will use default of {publicAccess=true, privateAccess=fal
se} for cluster "lwcluster" in "ap-south-1"
[0] 2 sequential tasks: { create cluster control plane "lwcluster", 2 sequential sub-tasks:
{ no tasks, 2 parallel sub-tasks: { create nodegroup "ng1", create nodegroup "ng2" } } }
[0] building cluster stack "eksctl-lwcluster-cluster"
[0] deploying stack "eksctl-lwcluster-cluster"

```

```

will create a CloudFormation stack for cluster itself
will create a CloudFormation stack for cluster itself

```



And this is doing by the EKSCTL automatically and this automatically doing is called as the Orchestration
And for this we have the service in AWS called as Cloud Formation
And the EKSCTL is internally creating a Cloud Formation program and this is the one that id doing the stuff for us

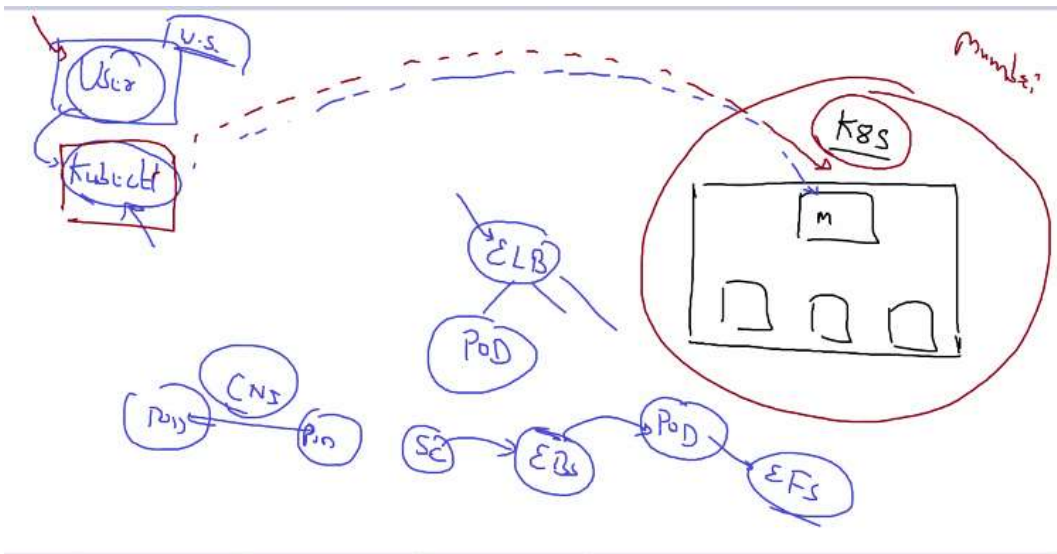
And here the Cloud Formation creation is called as Stack in Cloud

And after the cluster is created then we have to use it (and from here our stuff starts)

And for using we have many other things come in between

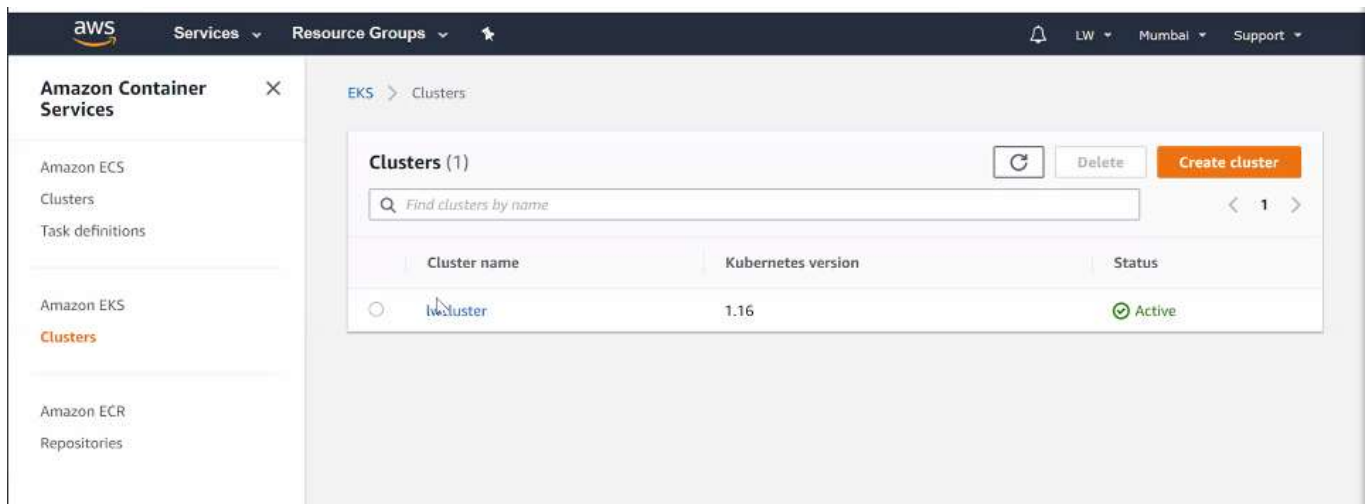
And using means we are using the K8s in EKS

And to use the CLI of the K8s we need to have the Kubectl program



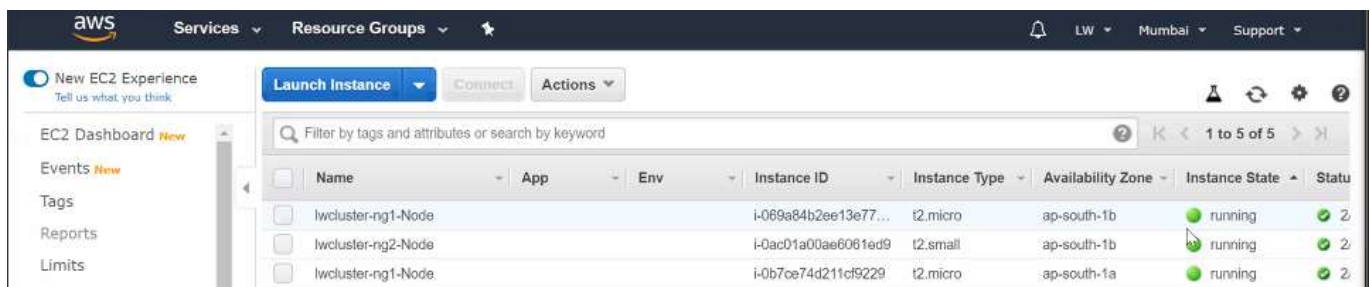
And we will be doing tons of integration and all this cannot be done in our laptop and can only be done using the AWS cloud or some other Cloud

And also the AWS is giving the feature of using its services as this is the one that is managing the



```
C:\Users\Vimal Daga\Desktop\eks_class_code>eksctl get cluster
NAME      REGION
lwcluster ap-south-1
```

now the cluster is created successfully



And in the EC2 we can see that we have 3 more nodes we had created

And here we have forgot to attach the Key pair

So without that we can't login into the node and see what is happening internally

```

apiVersion: eksctl.io/v1alpha5
kind: ClusterConfig

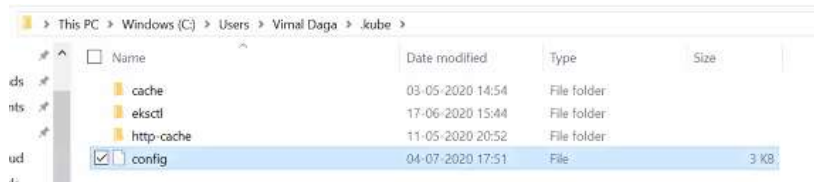
metadata:
  name: lwcluster
  region: ap-south-1

nodeGroups:
  - name: ng1
    desiredCapacity: 2
    instanceType: t2.micro
    ssh:
      publicKeyName: mykey111222
  - name: ng2
    desiredCapacity: 1
    instanceType: t2.small

```

And internally if we want to attach the key we can use the following code in the file

And in the Kubectl which we have to use for the



And in our case we don't find this file

And we have to create it

And in this we just write the IP of our cluster and then username and password (that's it)

```

C:\Users\Vimal Daga>kubectl config view
apiVersion: v1
clusters: null
contexts: null
current-context: ""
kind: Config
preferences: {}
users: null

```

And if we don't have that file then if we run this command then we know about the details in the file

And we have an aws command that will create this file for us

```

C:\Users\Vimal Daga>aws eks update-kubeconfig

```

And for this we have to pass which cluster we have to update this file

```

C:\Users\Vimal Daga>aws eks update-kubeconfig --name lwcluster
Added new context arn:aws:eks:ap-south-1:417149810339:cluster/lwcluster to C:\Users\Vimal Daga\.kube\config

```

And the contents in the file will be like


```

apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSB0tLS0tCk1JSUN5RENDQWJDZl
    dmI5VXorZUdTVETrcmIxQjBNS0l3TzFXd1ovNS8wNmJpTnpkbGV5RFNRC2NTTEFGTGf1ZjVXVd0Q1FmMEltCj
    server: https://7BC589ADB090295994F768EF134BBFF4.sk1.ap-south-1.eks.amazonaws.com
    name: arn:aws:eks:ap-south-1:417149810339:cluster/lwcluster
contexts:
- context:
    cluster: arn:aws:eks:ap-south-1:417149810339:cluster/lwcluster
    user: arn:aws:eks:ap-south-1:417149810339:cluster/lwcluster
    name: arn:aws:eks:ap-south-1:417149810339:cluster/lwcluster
current-context: arn:aws:eks:ap-south-1:417149810339:cluster/lwcluster
kind: Config
preferences: {}
users:
- name: arn:aws:eks:ap-south-1:417149810339:cluster/lwcluster
  user:
    exec:

```

And this is the IP or URL for the cluster we have

And then we can use the K8s service now

```

C:\Users\Vimal Daga\Desktop\eks_class_code>
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pods
No resources found in default namespace.

```

```

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get nodes
NAME                                STATUS    ROLES    AGE   VERSION
ip-192-168-57-189.ap-south-1.compute.internal Ready    <none>   21m   v1.16.8-eks-fd1ea7
ip-192-168-7-78.ap-south-1.compute.internal Ready    <none>   21m   v1.16.8-eks-fd1ea7
ip-192-168-9-210.ap-south-1.compute.internal Ready    <none>   20m   v1.16.8-eks-fd1ea7

```

And these are the nodes that are launched and also these are the IP's of the nodes

And these are the actually the EC2 IP's

And if we want to see some extra information about the node we use the describe command

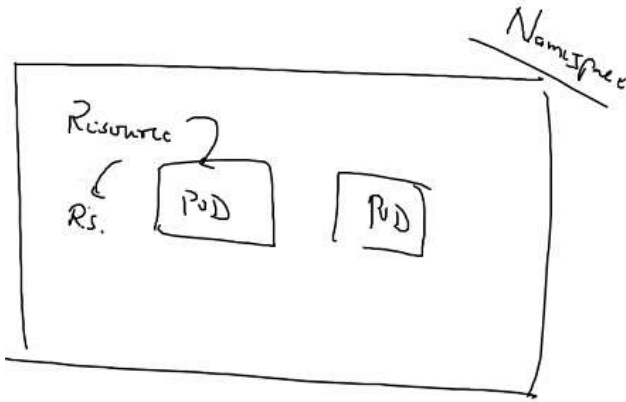
```

Container Runtime Version: docker://19.3.6
Kubelet Version: v1.16.8-eks-fd1ea7
Kube-Proxy Version: v1.16.8-eks-fd1ea7
ProviderID: aws:///ap-south-1b/i-069a84b2ee13e771a
Non-terminated Pods: (3 in total)
Namespace           Name                                CPU Requests  CPU Limits  Memory Requests  Memory Lim
its  AGE
-----
kube-system          aws-node-j8w9f                     10m (1%)     0 (0%)      0 (0%)          0 (0%)
22m
kube-system          coredns-6856799b8d-zqxzn          100m (10%)   0 (0%)      70Mi (11%)      170Mi (27%)
28m
kube-system          kube-proxy-ck7pc                   100m (10%)   0 (0%)      0 (0%)          0 (0%)
22m
Allocated resources:
(Total limits may be over 100 percent, i.e., overcommitted.)
Resource           Requests  Limits
-----
cpu                 210m (22%)  0 (0%)
memory              70Mi (11%)  170Mi (27%)
ephemeral-storage   0 (0%)      0 (0%)
hugepages-2Mi       0 (0%)      0 (0%)
attachable-volumes-aws-efs 0            0
Events:
Type      Reason      Age      From      Message

```

And here we see lots of information

And also if we have the PODS in K8s we keep all the related one in some box called as namespace



And we can see the namespace (and many are created by the K8s for us) and by default we are using the default Namespace

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get ns
NAME                STATUS   AGE
default             Active   32m
kube-node-lease     Active   32m
kube-public         Active   32m
kube-system         Active   32m
```

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl create namespace lwns
namespace/lwns created

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get ns
NAME                STATUS   AGE
default             Active   33m
kube-node-lease     Active   33m
kube-public         Active   33m
kube-system         Active   33m
lwns                Active   2s
```

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pods
No resources found in default namespace.
```

And if we want to change our default namespace then we will use the config file and then if we want to change the namespace then we have to add the context

```
>kubectl config view
```

This is another way to view the file

And we can change the default NS using the following command

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl config set-context --current --namespace=
lwns
Context "arn:aws:eks:ap-south-1:417149810339:cluster/lwcluster" modified.
```

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pods
No resources found in lwns namespace.
```

And till here is just the set up part and also about the use cases how the EKS is using the K8s

And now the fun part starts where we integrate features of services of AWS with the K8s (in EKS)

And if we want to see the cluster information then we can use the following command

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl cluster-info
Kubernetes master is running at https://7BC589ADB090295994F768EF1348BFF4.sk1.ap-south-1.eks.amazonaws.com
CoreDNS is running at https://7BC589ADB090295994F768EF1348BFF4.sk1.ap-south-1.eks.amazonaws.com/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
```

After the Break:

And now we have our cluster ready and now we can check for the practical and integration we want to do

And here we use some image that has the apache webserver configured like



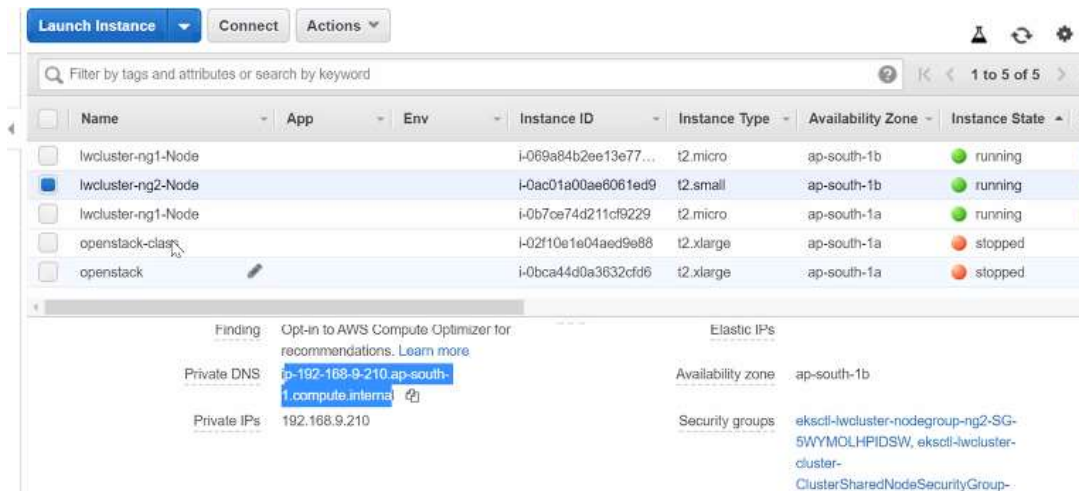
And here we can launch the deployment

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl create deployment myweb --image=vimal13/apache-webserver-php
deployment.apps/myweb created
```

And this way we can use this

And internally they are connected to the AWS and then EKS and then they are we have launched our POD in one of the node in the EC2 of the Amazon

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE
myweb-79b48fb9f5-c7wbn             1/1     Running   0           49s   192.168.2.156   ip-192-168-9-210.ap-south-1.compute-1.internal
```



And we can see that it is only one container in the POD as we have asked for one

And now if we want to increase it we can use the scaling

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl scale deployment myweb --replicas=3
deployment.apps/myweb scaled
```

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
myweb-79b48fb9f5-c7wbn             1/1    Running   0           2m29s
myweb-79b48fb9f5-p99hx             1/1    Running   0           7s
myweb-79b48fb9f5-xm5tc             0/1    ContainerCreating   0           7s
```

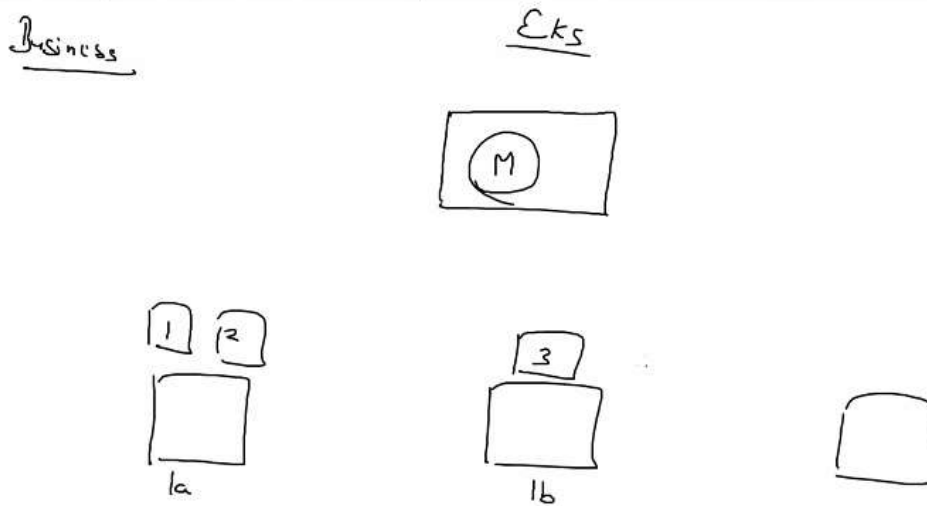
```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE
myweb-79b48fb9f5-c7wbn             1/1    Running   0           2m38s  192.168.2.156   ip-192-168-9-210.ap-south-1.com
ute.internal                       <none>   <none>     0           16s    192.168.11.0    ip-192-168-9-210.ap-south-1.com
myweb-79b48fb9f5-p99hx             1/1    Running   0           16s    192.168.11.0    ip-192-168-9-210.ap-south-1.com
ute.internal                       <none>   <none>     0           16s    192.168.34.113  ip-192-168-57-189.ap-south-1.com
myweb-79b48fb9f5-xm5tc             1/1    Running   0           16s    192.168.34.113  ip-192-168-57-189.ap-south-1.com
pute.internal                       <none>   <none>     0           16s    192.168.34.113  ip-192-168-57-189.ap-south-1.com
```

And the beauty of the cluster is that these pods are launched in different Nodes

And the one who plans this is the Scheduler

And the beauty here is that one node is in the different data centre and others are in other data centre

The present setup is like



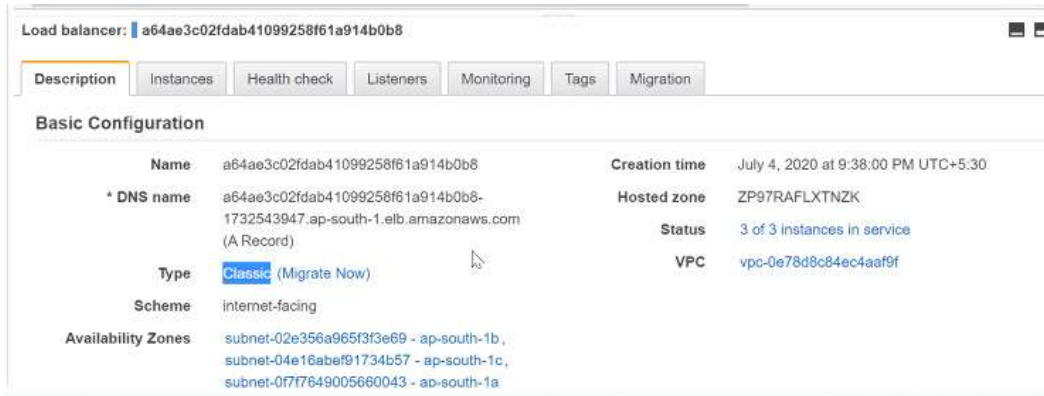
And here they are running all the nodes and instances in their own network or VPC and inside a Subnet

And for this we can't connect them from outside world as they are using the Public IP or in the Public World

And for this we have to create a program called Service or Load balancer

- And in K8s we have 3 types of these like
 - o Node Port (where we expose it using the Node IP)
 - o Cluster IP
 - o Load Balancer
- And here we will be using the Load Balancer service
- And this one give us the access to the outside world and also perform the load balancing
- And for this we have 2 types
 - o Either we can use the Node Port or
 - o Can use Load Balancer
- And today we will be using the Load Balancer

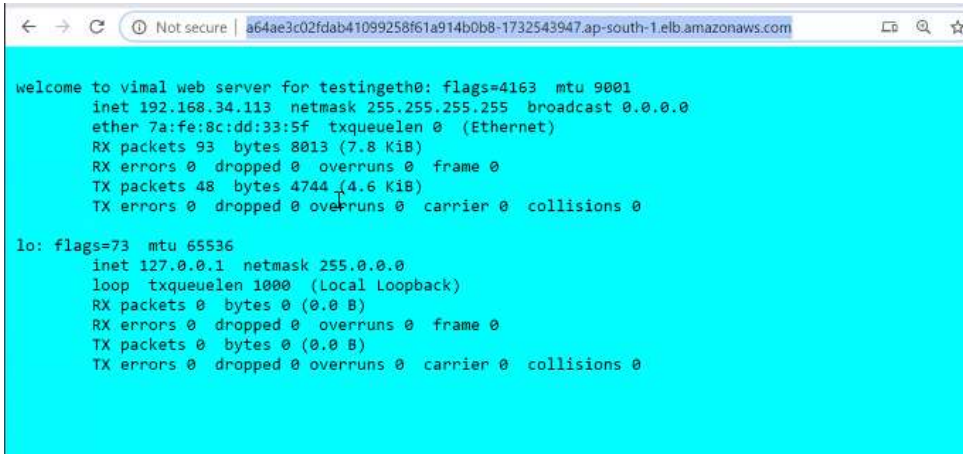
And the benefit of the ELB is that it will provide us the Public IP through which we can access the POD or instance from anywhere from the world



And in the ELB they are using the Classic Load Balancer and actually we have 3 types of the Load Balancer

And this the use case of the Load balancer with that of the Node Port in K8s

As when we want to connect some load Balancer then we can use the Load Balancer instead of Node Port



And now we can connect to the page or apache server

And here we have 2 load balancer

- One is the EKS load balancer
- And the other is the K8s Load Balancer

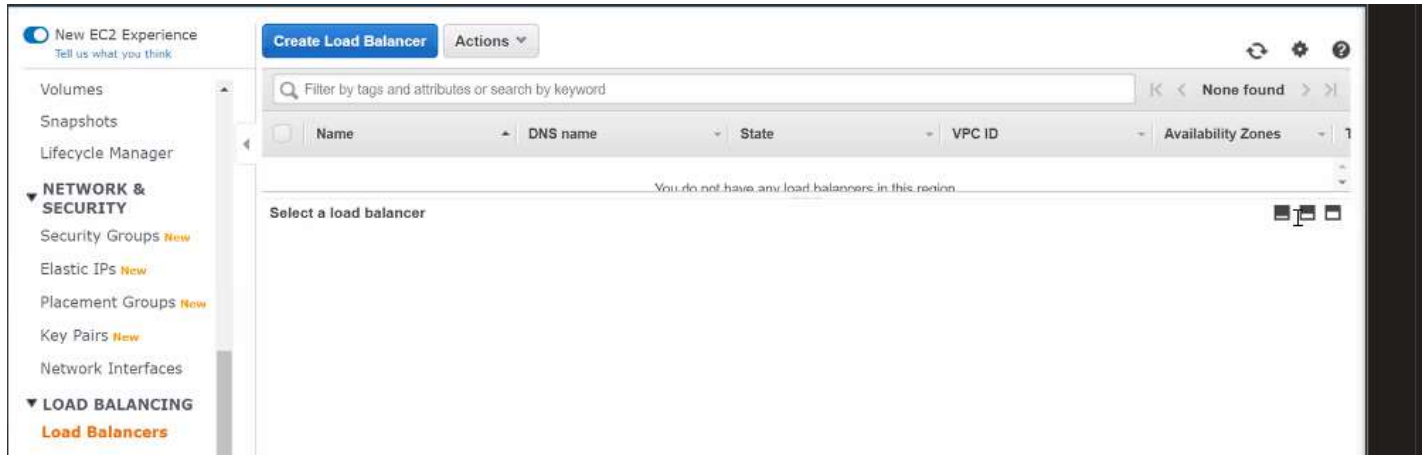
And here they are balancing it in between the PODS and also in between the Nodes and the nodes are in between different Data centres so we have the load balanced between the Data Centres too

It is like the Base Infrastructure is maintained by the AWS and the Management of the POD or application is done by the Instance

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl describe service/myweb
Name: myweb
Namespace: lwns
Labels: app=myweb
Annotations: <none>
Selector: app=myweb
Type: LoadBalancer
IP: 10.100.147.38
LoadBalancer Ingress: a64ae3c02fdab41099258f61a914b0b8-1732543947.ap-south-1.elb.amazonaws.com
Port: <unset> 80/TCP
TargetPort: 80/TCP
NodePort: <unset> 32753/TCP
Endpoints: 192.168.11.0:80,192.168.2.156:80,192.168.34.113:80
Session Affinity: None
External Traffic Policy: Cluster
Events:
```

And if we describe our Load Balancer or service then we observe that the Ingress (the input or the Hits) is received by this IP or URL

And now if we delete all in K8s then it will delete the service means the Load Balancer is deleted which means the ELB is also deleted and this we can observe this



And here the complete focus we can look into the App and the management of the resources are also done by the EKS

And if we want to go inside the container then we can use

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl exec -it myweb-79b48fb9f5-6hf46 bash
```

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl exec -it myweb-79b48fb9f5-6hf46 bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl kubectl exec [POD] -- [COMMAND] instead.
[root@myweb-79b48fb9f5-6hf46 /]#
```

And now we are inside the POD

- This is like
 - IN AWS cloud in Mumbai DC
 - We have a Cluster
 - In which we have Node (EC2)
 - And then to a POD in it
 - And we have logged into it from our home

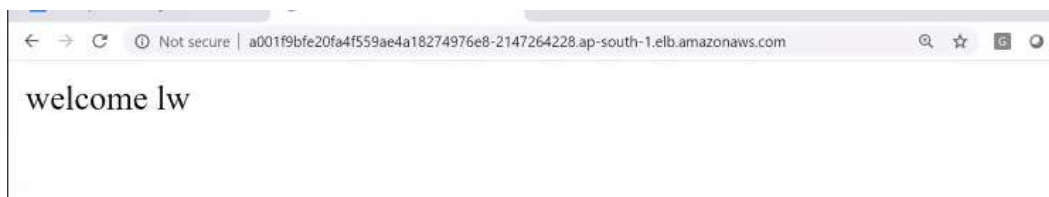
And say if we want to add or change something in the POD then we have to login into it and then we can change it

And we can do this without logging in to the POD

And for this we are just using the K8s command and we don't need to know anything about the AWS

And this can be done using the Copy Command

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl cp index.php myweb-79b48fb9f5-6hf46:/var/www/html/index.php
```



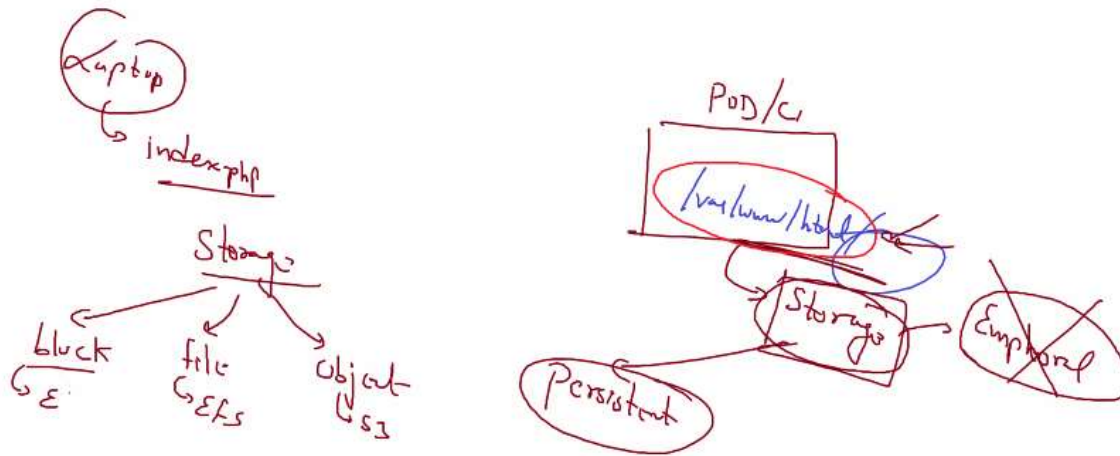
And this has changed the page

And whatever we have copied inside the POD then it is copied it into the POD storage and by default it is Ephemeral in nature (temporary)

And for this we need to store then or make it persistent

Say for this if we delete all the POD then it gets deleted and if we want to relaunch it automatically then we use the RS which is a controller and this will monitor the POD

And now it will relaunch it



But the new POD it made from the image so it will not have the data that we have copied in the Previous Node

And to make it persistent we have to attach a new storage to it

And in AWS it is called as Volume

And for this we can use 3 types of storage like

- Block (EBS)
- File (EFS)
- Object (S3)

And since we are connect a hard disk we have to use the Block Storage and also the then format and mount it

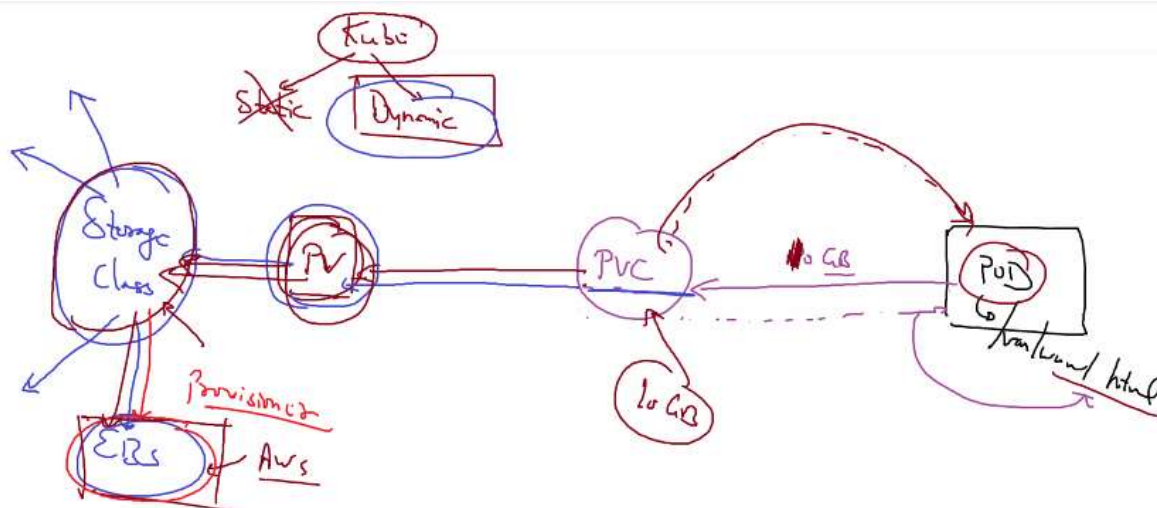
And this volume here in K8s (as we just want to use the EKS K8s to manage the resources) it is called as Persistent Volume

And the since we are asking for the PV it called as PVC

So the POD will ask for the PVC and the PVC will ask the PV for some space

and creating this PV we have 2 ways:

- Static
- Dynamic



And generally we are using the dynamic

And if we create a PVC then this will automatically that will create the PV

And the PV is asking from some program

As which program is the one that will manage from where we are getting the Physical Storage

And this program is called as Storage Class

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pvc
No resources found in lwns namespace.

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pv
No resources found in lwns namespace.

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get sc
NAME                PROVISIONER          AGE
gp2 (default)       kubernetes.io/aws-ebs 111m
```

And we have one storage class already created for us

```
NAME                PROVISIONER          AGE
gp2 (default)       kubernetes.io/aws-ebs 111m
```

And this storage class will use the EBS or contact to it if anybody asks it for the storage it will automatically give the storage and the volume in the AWS

And this is how they are connected internally the EKS and the AWS services

And for this we don't have any command but we have to use the code part using the YAML file code

And the file looks like somewhat of

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: lwpvc1
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
```

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl create -f pvc.yml
persistentvolumeclaim/lwpvc1 created
```

And since we have created the Storage Class we don't need to create the PV and will be created automatically when we create a PVC

But now if we see that the PV and hard disk is not created

Because the storage class is the one who will do that

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl describe sc gp2
Name:                gp2
IsDefaultClass:      Yes
Annotations:         kubectl.kubernetes.io/last-applied-configuration={"apiVersion":"storage.k8s.io/v1","kind":"StorageClass","metadata":{"annotations":{"storageclass.kubernetes.io/is-default-class":"true"},"name":"gp2"},"parameters":{"fsType":"ext4","type":"gp2"},"provisioner":"kubernetes.io/aws-ebs","volumeBindingMode":"WaitForFirstConsumer"}}
,storageclass.kubernetes.io/is-default-class=true
Provisioner:         kubernetes.io/aws-ebs
Parameters:          fsType=ext4,type=gp2
AllowVolumeExpansion: <unset>
MountOptions:         <none>
ReclaimPolicy:        Delete
VolumeBindingMode:    WaitForFirstConsumer
Events:
```

And if we describe that in the Annotation we have the Volume Bind Mode set to the Wait for the first Consume Means it will wait till the consumer is asked or uses the storage until then it will not create (for the 1st time)

And that's why we have the PVC pending


```
Events: <none>

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pvc
NAME      STATUS    VOLUME      CAPACITY   ACCESS MODES   STORAGECLASS   AGE
lwpvc1    Pending                10Gi         RWO         gp2             2m16s
```

But as soon as the consumer the POD will ask for the storage it will create the PVC and (it will create the PV automatically as the storage class is Specified)

And this can be set in the deployment file (as it is better to set in the deployment instead of POD)

```
kubectl-edit-px7s8 - Notepad
File Edit Format View Help

labels:
  app: myweb
spec:
  volumes:
    - name: web-vol1
      persistentVolumeClaim:
        claimName: lwpvc1
  containers:
    - image: vimal13/apache-webserver-php
      imagePullPolicy: Always
      name: apache-webserver-php
      resources: {}
      terminationMessagePath: /dev/termination-log
      terminationMessagePolicy: File
  dnsPolicy: ClusterFirst
  restartPolicy: Always
  schedulerName: default-scheduler
  securityContext: {}
  containers:
    - image: vimal13/apache-webserver-php
      volumeMounts:
        - mountPath: /var/www/html
          name: web-vol1
```

And now we have to tell the POD that it has to use the PVC

But it is better to say it to the deployment instead of the POD or container

And also tell where to mount in the POD

And this can be changed in the YAML file of the Deployment

```
C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl edit deployment.apps/myweb
deployment.apps/myweb edited

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
myweb-79b48fb9f5-xvqdc             1/1     Running   0           21m
myweb-7b7c54778f-rfwwd             0/1     ContainerCreating   0           7s

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pvc
NAME      STATUS    VOLUME                                     CAPACITY   ACCESS MODES   STORAGECLASS   AGE
lwpvc1    Bound     pvc-a7529549-6cbd-44a5-80e3-db758969e287  10Gi         RWO             gp2             7m48s
```

And now if we deploy it

Then the PODS are restarted as they are to be make the changes effective

And then we see that the PVC is also created and also if we describe the POD we see that the mount is also done


```

9274ca44118d38f4601c0080a91
Port:          <none>
Host Port:     <none>
State:         Running
  Started:     Sat, 04 Jul 2020 22:18:23 +0530
Ready:         True
Restart Count: 0
Environment:   <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-882tg (ro)
  /var/www/html from web-vol1 (rw)

```

And as we have mounted nothing is there in the mount point at the start by
But as soon as we copy some files then those are added to the storage

```

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl delete pods --all
pod "myweb-7b7c54778f-rfwwd" deleted

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
myweb-7b7c54778f-vrwt8             1/1     Running   0           15s

```

And now if we have deleted the pod still the data is not lost made persistent

And also when we create the storage in the AWS we have different types

Create Volume



And when we create the PVC we want that the K8s must use the IOPS instead of GP2 (this is by default)

```

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl describe sc
Name:          gp2
IsDefaultClass: Yes
Annotations:   kubectl.kubernetes.io/last-applied-configuration={"apiVersion":"storage.k8s.io/v1","kind":"StorageClass","metadata":{"annotations":{"storageclass.kubernetes.io/is-default-class":"true"},"name":"gp2"},"parameters":{"fsType":"ext4","type":"gp2"},"provisioner":"kubernetes.io/aws-ebs","volumeBindingMode":"WaitForFirstConsumer"}}
,storageclass.kubernetes.io/is-default-class=true
Provisioner:   kubernetes.io/aws-ebs
Parameters:    fsType=ext4,type=gp2
AllowVolumeExpansion: <unset>
MountOptions:   <none>
ReclaimPolicy:  Delete
VolumeBindingMode: WaitForFirstConsumer
Events:        <none>

```

And for this we have to create an another storage type and for this we have to change the type

And for this we have to create a storage class using the YAML file

```

apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: lwsc1
provisioner: kubernetes.io/aws-ebs
parameters:
  type: io1
reclaimPolicy: Retain

```

And also if we observe in the PV or Storage Class then we see that the ReClaim Policy is Delete

```

C:\Users\Vimal Daga\Desktop\eks_class_code>kubect1 get pv
NAME                                CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  CLAIM          ST
ORAGECLASS  REASON  AGE
pvc-a7529549-6cbd-44a5-80e3-db758969e287  10Gi      RWO           Delete          Bound   lwns/lwpvc1    gp
2
10m

```

Means if we delete the PVC then the HDD or Volume is also deleted

And if we want to retain the Volume instead of delete then we have to change the Re Claim Policy to Retain

```

C:\Users\Vimal Daga\Desktop\eks_class_code>kubect1 delete pvc --all
persistentvolumeclaim "lwpvc1" deleted

C:\Users\Vimal Daga\Desktop\eks_class_code>kubect1 get pvc
No resources found in lwns namespace.

C:\Users\Vimal Daga\Desktop\eks_class_code>kubect1 get pv
NAME                                CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  CLAIM          ST
ORAGECLASS  REASON  AGE
pvc-a7529549-6cbd-44a5-80e3-db758969e287  10Gi      RWO           Delete          Failed   lwns/lwpvc1    gp
2
12m

C:\Users\Vimal Daga\Desktop\eks_class_code>kubect1 get pv
NAME                                CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  CLAIM          ST
ORAGECLASS  REASON  AGE
pvc-a7529549-6cbd-44a5-80e3-db758969e287  10Gi      RWO           Delete          Failed   lwns/lwpvc1    gp
2
12m

C:\Users\Vimal Daga\Desktop\eks_class_code>kubect1 get pv
No resources found in lwns namespace.

```

And now we have the 2 Storage Classes

```

C:\Users\Vimal Daga\Desktop\eks_class_code>kubect1 create -f sc.yml
storageclass.storage.k8s.io/lwsc1 created

C:\Users\Vimal Daga\Desktop\eks_class_code>kubect1 get sc
NAME                PROVISIONER  AGE
gp2 (default)      kubernetes.io/aws-ebs  138m
lwsc1               kubernetes.io/aws-ebs  6s

```

And in the PVC we never written the Storage class to be used as it will pick the by default Storage Class

But if we want to use the one we have created then we have to add the Storage class Section to tell which on to be used

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: lwpvc1
spec:
  storageClassName: lwsc1
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi

```

```

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pvc
NAME      STATUS    VOLUME                                     CAPACITY   ACCESS MODES   STORAGECLASS   AGE
lwpvc1    Pending                                     10Gi        RWO          lwsc1          4s

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pv
NAME      STATUS    RECLAIM POLICY   CAPACITY   ACCESS MODES   STORAGECLASS   AGE
pvc-a7b85e01-7783-401c-8547-e51dcf018fc3  10Gi        RWO          Retain      Bound          lwns/lwpvc1    1w
sc1
5s

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pvc
NAME      STATUS    VOLUME                                     CAPACITY   ACCESS MODES   STORAGECLASS   AGE
lwpvc1    Bound     pvc-a7b85e01-7783-401c-8547-e51dcf018fc3  10Gi        RWO          lwsc1          16s

```

And now here we have created the volume without using or asked by the consumer

And now if we describe this Storage Class then

- We have not written any annotations but was there in the previous one

```

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get sc
NAME      PROVISIONER   AGE
gp2 (default)  kubernetes.io/aws-ebs  143m
lwsc1 (default) kubernetes.io/aws-ebs  4m52s

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl edit sc gp2
storageclass.storage.k8s.io/gp2 edited

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl edit sc gp2
Edit cancelled, no changes made.

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get sc
NAME      PROVISIONER   AGE
gp2 (default)  kubernetes.io/aws-ebs  144m
lwsc1 (default) kubernetes.io/aws-ebs  5m29s

```

And if we see that both are default now but this raises to the dilemma so

Then we have to change it in the other one by setting the is-default-class to false

```
storageclass.kubernetes.io/is-default-class: "false"
```

And now if we delete the PVC then we observe that the

```

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl delete pvc --all
persistentvolumeclaim "lwpvc1" deleted

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pvc
No resources found in lwns namespace.

C:\Users\Vimal Daga\Desktop\eks_class_code>kubectl get pv
NAME      STATUS    RECLAIM POLICY   CAPACITY   ACCESS MODES   STORAGECLASS   AGE
pvc-a7b85e01-7783-401c-8547-e51dcf018fc3  10Gi        RWO          Retain      Released          lwns/lwpvc1
lwsc1
4m23s

```

The PV is retained

```
C:\Users\Vimal Daga\Desktop\eks_class_code>eksctl delete cluster -f cluster.yml .
```

And if we use this command then the entire cluster is also will be deleted in one go

As internally it will delete the Cloud Formation template and as this is deleted the entire cluster is also deleted

And if we want to use these commands then we have to have the admin power so for this we have to create a user